

ACTIVITY 3.1: -

Question-1: - Find genres of books including duplicates

Answer: - select genre from books;

```
spr2022adb35=> select genre from books;
      genre
-----
Horror
Tragedy
Autobiographical
Drama/Romance
Drama
Romance/Tragedy
Romance/Diary
Travel
Historical/Drama
Historical/Drama
Romance
Historical
Fantasy/Adventure
Romance
Romance
(15 rows)
```

Question-2: - Find genres of books excluding duplicates

Answer: - select distinct(genre) from books;

```
spr2022adb35=> select distinct(genre) from books;
      genre
-----
Romance
Fantasy/Adventure
Drama/Romance
Autobiographical
Travel
Historical
Drama
Horror
Romance/Diary
Historical/Drama
Tragedy
Romance/Tragedy
(12 rows)

spr2022adb35=>
```

Question-3: - Find all horror books with alias booktitle for title.

Answer: - select id, title as booktitle, pagecount,genre,authorid,pubid from books where genre='Horror';

```
OpenSSH SSH client
spr2022adb35=> select id, title as booktitle, pagecount,genre,authorid,pubid from books where genre='Horror';
id | booktitle | pagecount | genre | authorid | pubid
-----+-----
1 | It        | 1138      | Horror | 10        | 100
(1 row)
```

```
spr2022adb35=> select title as booktitle, genre from books where genre='Horror';
booktitle | genre
-----+-----
It        | Horror
(1 row)
```

```
spr2022adb35=> select id, title as booktitle from books where genre='Horror';
id | booktitle
-----+-----
1 | It
(1 row)
```

Question-4: - Find the languages of all editions of “it”.


Answer: - select distinct lang, edition, title from books,editions where books.id = editions.bookid and title='It';

```
OpenSSH SSH client
spr2022adb35=> select distinct lang, edition, title from books,editions where books.id = editions.bookid and title='It';
lang | edition | title
-----+-----+-----
English | It (Hardcover) | It
German | Es        | It
Italian | It (Paperback) | It
(3 rows)
```

ACTIVITY-3.2(a) Aggregates: - Try a few queries

Question-1: - Try the select query

Answer: - select max(pagecount) from books;

 OpenSSH SSH client

```
spr2022adb35=> select max(pagecount) from books;
max
-----
1349
(1 row)
```

Question-2: - Adding where clause to question-1.

Answer: - select max(pagecount) from books where genre = 'Drama'

```
spr2022adb35=> select max(pagecount) from books where genre = 'Drama'
spr2022adb35-> ;
max
-----
288
(1 row)
```

Question-3: - select query with error

Answer: - select title, max(pagecount) from books where genre='Drama';

It gives an error message as given below: -

```
spr2022adb35=> select title, max(pagecount) from books where
genre='Drama';
```

ERROR: column "books.title" must appear in the GROUP BY clause or be used in an aggregate function

LINE 1: select title, max(pagecount) from books where genre='Drama';

```
spr2022adb35=> select title, max(pagecount) from books where genre='Drama';
ERROR: column "books.title" must appear in the GROUP BY clause or be used in an aggregate function
LINE 1: select title, max(pagecount) from books where genre='Drama';
                  ^
```

This error can be resolved by adding the group by clause as given below: -

`select title, max(pagecount) from books where genre='Drama' group by title;`

```
spr2022adb35=> select title, max(pagecount) from books where genre='Drama' group by title;
      title      | max
-----+-----
The Joy Luck Club | 288
(1 row)
```

Activity 3.2(b) Aggregates: Average and Max

Question-1: - Example instance for different averages

Answer: - Both of these queries will yield different results when there is a duplicate entry for pagecount.

In this case, there is a duplicate value as highlighted in the screenshot below: -

Pagecount-224 for book titles September love and The nickel boys.

This means when distinct pagecount was not used for calculating the average of pagecounts, then the result was 432.4667(approx.) whereas when distinct pagecount was used, then the result was 447.357 (approx.).

This is also present in the screenshot presented on the next page.

```
spr2022adb35=>
spr2022adb35=> select pagecount,title from books;
pagecount |
-----+-----
1138 | It
500 | Hamlet
304 | I Know Why the Caged Bird Sings
1349 | A Suitable Boy
288 | The Joy Luck Club
256 | Like Water for Chocolate
294 | Tita's Diary
464 | From Heaven Lake
371 | Kite Runner
352 | The Vanishing Half
224 | September Love
224 | The Nickel Boys
163 | The Alchemist
176 | Love and Misadventure
384 | The Authenticity Project
(15 rows)

spr2022adb35=> select distinct(pagecount),title from books;
pagecount |
-----+-----
500 | Hamlet
256 | Like Water for Chocolate
352 | The Vanishing Half
224 | The Nickel Boys
304 | I Know Why the Caged Bird Sings
176 | Love and Misadventure
1349 | A Suitable Boy
224 | September Love
294 | Tita's Diary
371 | Kite Runner
288 | The Joy Luck Club
163 | The Alchemist
464 | From Heaven Lake
1138 | It
384 | The Authenticity Project
(15 rows)

spr2022adb35=>
spr2022adb35=> select avg(pagecount) from books;
avg
-----
432.466666666666667
(1 row)

spr2022adb35=> select avg(distinct pagecount) from books;
avg
-----
447.3571428571428571
(1 row)

spr2022adb35=>
```

Question-2: - can max(attribute) return different values or not?

Answer: - No, the value of these queries will always be the same and that is equal to the highest value of the attribute present in the table.

There is only one condition where the max(attribute) query can return different results (when distinct is used and when not) is only when a group by clause is used. This is because group by does a grouping according to the attribute and that can vary the result.

This can also be proved by using the screenshot below: -

 OpenSSH SSH client

```
spr2022adb35=> select max(pagecount) from books;
max
-----
1349
(1 row)

spr2022adb35=> select max(distinct(pagecount)) from books;
max
-----
1349
(1 row)
```

ACTIVITY 3.3(A): - INSERTS

Question-1: - Insert query for books

Answer: - insert into books values(16,'Paper Towns',420,'Young adult',24,100);

This is also highlighted in the screenshot below along with the select * to confirm the insertion.

```
Select OpenSSH SSH client
spr2022adb35=> \d books;
Table "spr2022adb35.books"
Column | Type | Collation | Nullable | Default
id      | integer |          | not null |
title   | text    |          |          |
pagecount | integer |          |          |
genre    | text    |          |          |
authorid | integer |          |          |
pubid    | integer |          |          |
Indexes:
    "books_pkey" PRIMARY KEY, btree (id)

spr2022adb35=>
spr2022adb35=> select * from books;
 id | title | pagecount | genre | authorid | pubid
-----|-----|-----|-----|-----|-----
 1 | It | 1138 | Horror | 10 | 100
 2 | Hamlet | 500 | Tragedy | 13 | 103
 3 | I Know Why the Caged Bird Sings | 304 | Autobiographical | 14 | 102
 4 | A Suitable Boy | 1349 | Drama/Romance | 15 | 103
 5 | The Joy Luck Club | 288 | Drama | 16 | 104
 6 | Like Water for Chocolate | 256 | Romance/Tragedy | 17 | 105
 7 | Tita's Diary | 294 | Romance/Diary | 17 | 105
 8 | From Heaven Lake | 464 | Travel | 15 | 102
 9 | Kite Runner | 371 | Historical/Drama | 18 | 106
10 | The Vanishing Half | 352 | Historical/Drama | 19 | 106
11 | September Love | 224 | Romance | 20 | 107
12 | The Nickel Boys | 224 | Historical | 21 | 108
13 | The Alchemist | 163 | Fantasy/Adventure | 22 | 103
14 | Love and Misadventure | 176 | Romance | 20 | 107
15 | The Authenticity Project | 384 | Romance | 23 | 102
(15 rows)

spr2022adb35=> insert into books values(16,'Paper Towns',420,'Young adult',24,100);
INSERT 0 1
spr2022adb35=> select * from books;
 id | title | pagecount | genre | authorid | pubid
-----|-----|-----|-----|-----|-----
 1 | It | 1138 | Horror | 10 | 100
 2 | Hamlet | 500 | Tragedy | 13 | 103
 3 | I Know Why the Caged Bird Sings | 304 | Autobiographical | 14 | 102
 4 | A Suitable Boy | 1349 | Drama/Romance | 15 | 103
 5 | The Joy Luck Club | 288 | Drama | 16 | 104
 6 | Like Water for Chocolate | 256 | Romance/Tragedy | 17 | 105
 7 | Tita's Diary | 294 | Romance/Diary | 17 | 105
 8 | From Heaven Lake | 464 | Travel | 15 | 102
 9 | Kite Runner | 371 | Historical/Drama | 18 | 106
10 | The Vanishing Half | 352 | Historical/Drama | 19 | 106
11 | September Love | 224 | Romance | 20 | 107
12 | The Nickel Boys | 224 | Historical | 21 | 108
13 | The Alchemist | 163 | Fantasy/Adventure | 22 | 103
14 | Love and Misadventure | 176 | Romance | 20 | 107
15 | The Authenticity Project | 384 | Romance | 23 | 102
16 | Paper Towns | 420 | Young adult | 24 | 100
(16 rows)
```


Question-2: - Insertion using the /COPY command

Answer: - The first step is to create the csv file on the local PC.

The file's contents are attached below along with the file: -

id	title	pagecount	genre	authorid	pubid
17	Looking for	620	Young Adu	24	100



books-copy.csv

After this, we must copy this file from local machine to the remote machine (linux.pdx.cs.edu). This will be done by the following command: -

scp books-copy.csv parth2@131.252.208.103:/u/parth2/DBMS/bookschema

After this, we will use the following command to copy the contents of the file into the database: -

\copy books from books-copy.csv with csv header

The output along with the select command highlighting the inserted field is given below: -

Select OpenSSH SSH client

```
spr2022adb35=> \copy books from books.csv with csv header
books-copy.csv books.csv
spr2022adb35=> \copy books from books-copy.csv with csv header
```

```
COPY 1
```

```
spr2022adb35=>
```

```
spr2022adb35=>
```

```
spr2022adb35=>
```

```
spr2022adb35=> select * from books;
```

id	title	pagecount	genre	authorid	pubid
1	It	1138	Horror	10	100
2	Hamlet	500	Tragedy	13	103
3	I Know Why the Caged Bird Sings	304	Autobiographical	14	102
4	A Suitable Boy	1349	Drama/Romance	15	103
5	The Joy Luck Club	288	Drama	16	104
6	Like Water for Chocolate	256	Romance/Tragedy	17	105
7	Tita's Diary	294	Romance/Diary	17	
8	From Heaven Lake	464	Travel	15	102
9	Kite Runner	371	Historical/Drama	18	106
10	The Vanishing Half	352	Historical/Drama	19	106
11	September Love	224	Romance	20	107
12	The Nickel Boys	224	Historical	21	108
13	The Alchemist	163	Fantasy/Adventure	22	103
14	Love and Misadventure	176	Romance	20	107
15	The Authenticity Project	384	Romance	23	102
16	Paper Towns	420	Young adult	24	100
17	Looking for Alaska	620	Young Adult	24	100

```
(17 rows)
```

```
spr2022adb35=>
```

Activity-3.3(B): -Default and null values

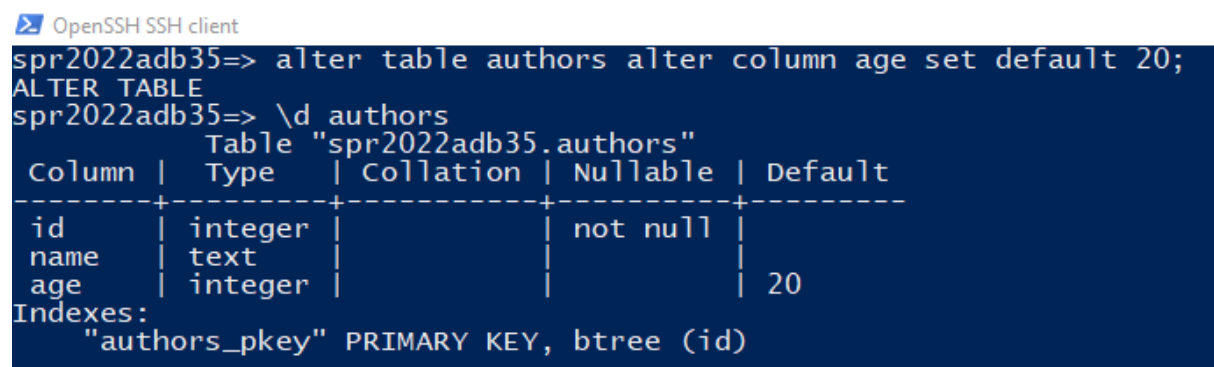
Question: - add default value to age in author, insert a author without an age and query to return author added above

Answer: - The **first task** to be done is to add a default value to the author's age in author table.

This will be done by the following query: -

alter table authors alter column age set default 20;

This change can be seen in the screenshot below: -



```
OpenSSH SSH client
spr2022adb35=> alter table authors alter column age set default 20;
ALTER TABLE
spr2022adb35=> \d authors
Table "spr2022adb35.authors"
Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----
id      | integer |          | not null |
name    | text    |          |          |
age     | integer |          |          | 20
Indexes:
    "authors_pkey" PRIMARY KEY, btree (id)
```

The **second task** to do is to insert an author without an age so that it sets to the default age. This will be done by the following query: -

insert into authors values(24, 'John Green');

To verify the insertion, we have the screenshot containing the highlighted insertion.

```
spr2022adb35=> insert into authors values(24, 'John Green');
INSERT 0 1
spr2022adb35=> select * from authors;
```

id	name	age
10	Stephen King	71
12	J.K. Rowling	53
13	Shakespeare	
14	Maya Angelou	
15	Vikram Seth	68
16	Amy Tan	69
17	Laura Esquivel	70
18	Khaled Hosseini	55
19	Brit Bennett	31
20	Lang Leav	40
21	Colson Whitehead	51
22	Paulo Coelho	73
23	clare pooley	49
100	parth parashar	50
24	John Green	20

```
(15 rows)
spr2022adb35=>
```

The **third task** to complete is to write a select query to return the author added in the second task: -

```
select * from authors where age=20;
```

```
spr2022adb35=> select * from authors where age=20;
```

id	name	age
24	John Green	20

```
(1 row)
```

ACTIVITY-3.4: - Updates and Null Values

Question-1: - Update the tuple inserted in 3.4 to set their age to null?

Answer: - update authors set age=NULL where id = 24;

The screenshot for this query is given below: -

```
OpenSSH SSH client
spr2022adb35=> update authors set age=NULL where id = 24;
UPDATE 1
spr2022adb35=>
spr2022adb35=>
```

Question-2: - Write a query to find authors where age is null?

Answer: - select * from authors where age is null;

The screenshot for the same is given below: -

```
spr2022adb35=>
spr2022adb35=>
spr2022adb35=> select * from authors where age is null;
 id |      name      | age
----+-----+-----
 13 | Shakespeare    |
 14 | Maya Angelou   |
 24 | John Green     |
(3 rows)
```

Question-3: - Write a query to find average age of authors? How do null values affect this average?

Answer: - select avg(age) from authors;

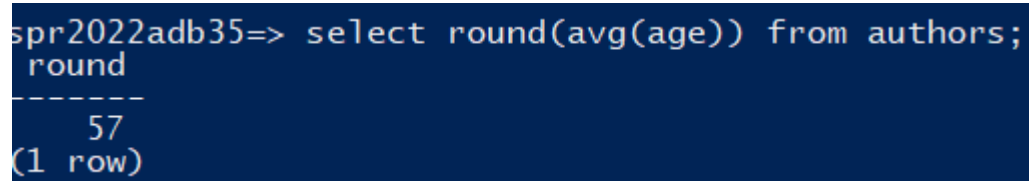
The result for this query is given in the screenshot below: -

```
OpenSSH SSH client
spr2022adb35=> select avg(age) from authors;
      avg
-----
56.666666666666667
(1 row)
```

We can also find the round value of the average age using the following query:

```
select round(avg(age)) from authors;
```

The screenshot for the same is given below: -




```
pr2022adb35=> select round(avg(age)) from authors;
round
-----
    57
(1 row)
```

Null values have no effect on the average age of authors. This is because NULL values are not included by the AVG aggregate function when calculating the average.

ACTIVITY-4.1: - RA – Select Activity

Question: - Find all books with pagecount>1000

Answer: -



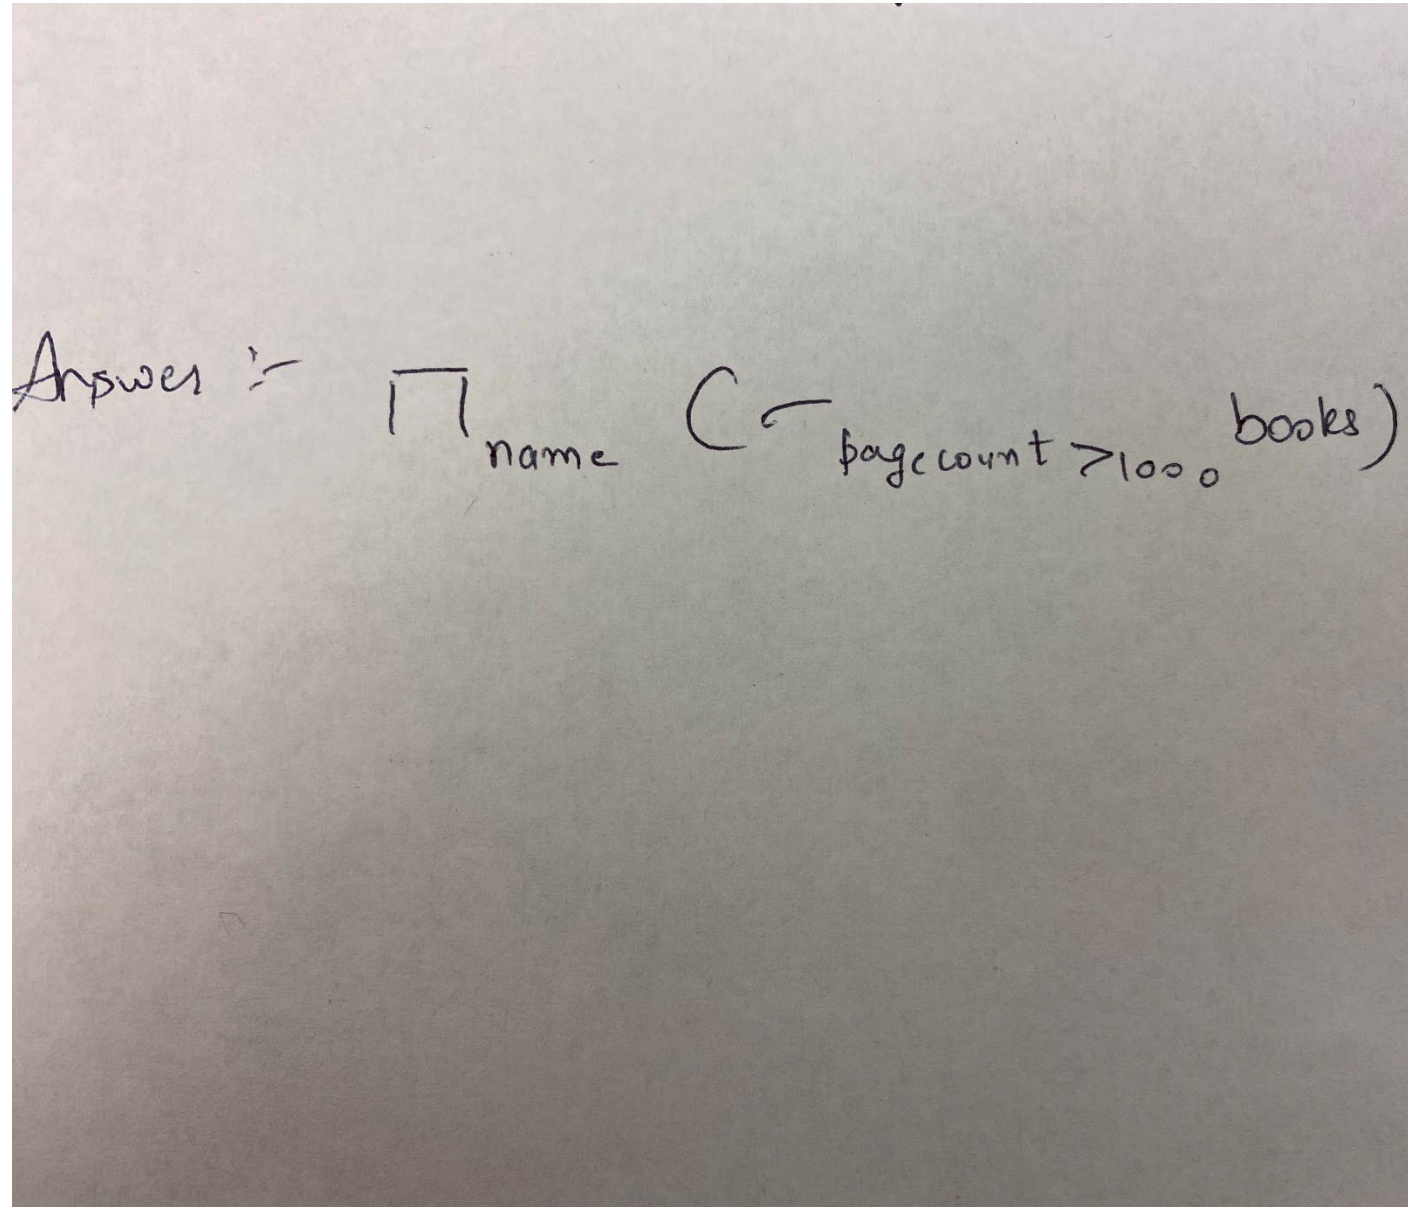
The equivalent SQL command is: -

Select * from books where pagecount>1000;

ACTIVITY-4.2 RA-Project Activity

Question-1: - Find all names of books with pagecount>1000.

Answer: -



The equivalent SQL statement is: -

Select name from books where pagecount>1000;

ACTIVITY-4.3-RA-Join Activity

Question: - Find the names of the author with books of pagecount>1000.

Answer: -

Answer:- $\pi_{\text{name}} (\sigma_{\text{pagecount} > 1000 \text{ and } \text{id} = \text{authorid}} (\text{authors} \times \text{books}))$

$\pi_{\text{name}} (\sigma_{\text{pagecount} > 1000} (\text{authors} \bowtie_{\text{id} = \text{authorid}} \text{books}))$

The equivalent SQL query is: -

select name from authors A JOIN books B on A.id = B.authorid where
pagecount>1000;

```
spr2022adb35=> select name from authors A JOIN books B on A.id = B.authorid where pagecount>1000;
name
-----
Stephen King
Vikram Seth
(2 rows)
```